

The Big LLM Architecture Comparison

From DeepSeek V3 to GLM-5: A Look At Modern LLM Architecture Design



SEBASTIAN RASCHKA, PHD

JUL 19, 2025

♡ 1,975

💬 97

🔄 174

Last updated: Apr 2, 2026 (added Gemma 4 in section 23)

It has been seven years since the original GPT architecture was developed. At first glar looking back at GPT-2 (2019) and forward to DeepSeek V3 and Llama 4 (2024-2025 might be surprised at how structurally similar these models still are.

Sure, positional embeddings have evolved from absolute to rotational (RoPE), Multi-He Attention has largely given way to Grouped-Query Attention, and the more efficient Sw has replaced activation functions like GELU. But beneath these minor refinements, hav truly seen groundbreaking changes, or are we simply polishing the same architectural foundations?

Comparing LLMs to determine the key ingredients that contribute to their good (or not good) performance is notoriously challenging: datasets, training techniques, and hyperparameters vary widely and are often not well documented.

However, I think that there is still a lot of value in examining the structural changes of th architectures themselves to see what LLM developers are up to in 2025. (A subset of t are shown in Figure 1 below.)




Figure 1: A subset of the architectures covered in this article.

So, in this article, rather than writing about benchmark performance or training algorithm will focus on the architectural developments that define today's flagship open models.

(As you may remember, [I wrote about multimodal LLMs](#) not too long ago; in this article, focus on the text capabilities of recent models and leave the discussion of multimodal capabilities for another time.)

Tip: This is a fairly comprehensive article, so I recommend using the navigation bar to go to the table of contents (just hover over the left side of the Substack page).

Optional: The video below is a narrated and abridged version of this article.


0:00

1. DeepSeek V3/R1

As you have probably heard more than once by now, [DeepSeek R1](#) made a big impact it was released in January 2025. DeepSeek R1 is a reasoning model built on top of the [DeepSeek V3 architecture](#), which was introduced in December 2024.

While my focus here is on architectures released in 2025, I think it's reasonable to include DeepSeek V3, since it only gained widespread attention and adoption following the launch of DeepSeek R1 in 2025.

If you are interested in the training of DeepSeek R1 specifically, you may also find my article from earlier this year useful:



Understanding Reasoning LLMs

SEBASTIAN RASCHKA, PHD • 2025年2月5日

[Read full story →](#)

In this section, I'll focus on two key architectural techniques introduced in DeepSeek V: improved its computational efficiency and distinguish it from many other LLMs:

- Multi-Head Latent Attention (MLA)
- Mixture-of-Experts (MoE)

1.1 Multi-Head Latent Attention (MLA)

Before discussing Multi-Head Latent Attention (MLA), let's briefly go over some background to motivate why it's used. For that, let's start with Grouped-Query Attention (GQA), which has become the new standard replacement for a more compute- and parameter-efficient alternative to Multi-Head Attention (MHA) in recent years.

So, here's a brief GQA summary. Unlike MHA, where each head also has its own set of keys and values, to reduce memory usage, GQA groups multiple heads to share the same key and value projections.

For example, as further illustrated in Figure 2 below, if there are 2 key-value groups and 4 attention heads, then heads 1 and 2 might share one set of keys and values, while heads 3 and 4 share another. This reduces the total number of key and value computations, which leads to lower memory usage and improved efficiency (without noticeably affecting the modeling performance, according to ablation studies).




Figure 2: A comparison between MHA and GQA. Here, the group size is 2, where a key and value pair is shared among 2 queries.

So, the core idea behind GQA is to reduce the number of key and value heads by sharing them across multiple query heads. This (1) lowers the model's parameter count and (2) reduces the memory bandwidth usage for key and value tensors during inference since keys and values need to be stored and retrieved from the KV cache.

(If you are curious how GQA looks in code, see my [GPT-2 to Llama 3 conversion guide](#) version without KV cache and my KV-cache variant [here](#).)

While GQA is mainly a computational-efficiency workaround for MHA, ablation studies as those in the [original GQA paper](#) and the [Llama 2 paper](#)) show it performs comparably to standard MHA in terms of LLM modeling performance.

Now, Multi-Head Latent Attention (MLA) offers a different memory-saving strategy that pairs particularly well with KV caching. Instead of sharing key and value heads like GQA, it compresses the key and value tensors into a lower-dimensional space before storing them in the KV cache.

At inference time, these compressed tensors are projected back to their original size before being used, as shown in the Figure 3 below. This adds an extra matrix multiplication but reduces memory usage.

DeepSeek V3/R1




Figure 3: Comparison between MLA (used in DeepSeek V3 and R1) and regular MHA.

(As a side note, the queries are also compressed, but only during training, not inference)

By the way, MLA is not new in DeepSeek V3, as its [DeepSeek-V2 predecessor](#) also used (and even introduced) it. Also, the V2 paper contains a few interesting ablation studies that may explain why the DeepSeek team chose MLA over GQA (see Figure 4 below).

unlike what previous papers found, this paper finds that MHA performs much BETTER than GQA

Benchmark (Metric)	# Shots	Dense 7B w/ MQA	Dense 7B w/ GQA (8 Groups)	Dense 7B w/ MHA
# Params	-	7.1B	6.9B	6.9B
BBH (EM)	3-shot	33.2	35.6	37.0
MMLU (Acc.)	5-shot	37.9	41.2	45.2
C-Eval (Acc.)	5-shot	30.0	37.7	42.9
CMMLU (Acc.)	5-shot	34.6	38.4	43.5

Table 8 | Comparison among 7B dense models with MHA, GQA, and MQA, respectively. MHA demonstrates significant advantages over GQA and MQA on hard benchmarks.

Benchmark (Metric)	# Shots	Small MoE w/ MHA	Small MoE w/ MLA	Large MoE w/ MHA	Large MoE w/ MLA
# Activated Params	-	2.5B	2.4B	25.0B	21.5B
# Total Params	-	15.8B	15.7B	250.8B	247.4B
KV Cache per Token (# Element)	-	110.6K	15.6K	860.2K	34.6K
BBH (EM)	3-shot	37.9	39.0	46.6	50.7
MMLU (Acc.)	5-shot	48.7	50.0	57.5	59.0
C-Eval (Acc.)	5-shot	51.6	50.9	57.9	59.2
CMMLU (Acc.)	5-shot	52.3	53.4	60.7	62.5

Table 9 | Comparison between MLA and MHA on hard benchmarks. DeepSeek-V2 shows better performance than MHA, but requires a significantly smaller amount of KV cache.

The memory requirements for MLA are much lower than for MHA

MLA performs better than MHA (here tested on Mixture-of-Experts architectures)

Figure 4: Annotated tables from the DeepSeek-V2 paper,
<https://arxiv.org/abs/2405.04434>

As shown in Figure 4 above, GQA appears to perform worse than MHA, whereas MLA better modeling performance than MHA, which is likely why the DeepSeek team chose over GQA. (It would have been interesting to see the "KV Cache per Token" savings comparison between MLA and GQA as well!)

To summarize this section before we move on to the next architecture component, ML, clever trick to reduce KV cache memory use while even slightly outperforming MHA in of modeling performance.

1.2 Mixture-of-Experts (MoE)

The other major architectural component in DeepSeek worth highlighting is its use of Mixture-of-Experts (MoE) layers. While DeepSeek did not invent MoE, it has seen a resurgence this year, and many of the architectures we will cover later also adopt it.

You are likely already familiar with MoE, but a quick recap may be helpful.

The core idea in MoE is to replace each FeedForward module in a transformer block with multiple expert layers, where each of these expert layers is also a FeedForward module. This means that we swap a single FeedForward block for multiple FeedForward blocks, as illustrated in the Figure 5 below.




Figure 5: An illustration of the Mixture-of-Experts (MoE) module in DeepSeek V3/R1 (right) compared to an LLM with a standard FeedForward block (left).

The FeedForward block inside a transformer block (shown as the dark gray block in the figure above) typically contains a large number of the model's total parameters. (Note that the transformer block, and thereby the FeedForward block, is repeated many times in a LLM; in the case of DeepSeek V3, 61 times.)

So, replacing *a single* FeedForward block with *multiple* FeedForward blocks (as done in MoE setup) substantially increases the model's total parameter count. However, the key is that we don't use ("activate") all experts for every token. Instead, a router selects only a small subset of experts per token. (In the interest of time, or rather article space, I'll cover the router in more detail another time.)

Because only a few experts are active at a time, MoE modules are often referred to as *sparse* in contrast to *dense* modules that always use the full parameter set. However, the large number of parameters via an MoE increases the capacity of the LLM, which means it can take up more knowledge during training. The sparsity keeps inference efficient, though we don't use all the parameters at the same time.

For example, DeepSeek V3 has 256 experts per MoE module and a total of 671 billion parameters. Yet during inference, only 9 experts are active at a time (1 shared expert plus 8 selected by the router). This means just 37 billion parameters are used per inference step opposed to all 671 billion.

One notable feature of DeepSeek V3's MoE design is the use of a shared expert. This is a shared expert that is always active for every token. This idea is not new and was already introduced in the [DeepSeek 2024 MoE](#) and [2022 DeepSpeedMoE papers](#).

Early MoE: Has bigger and fewer experts, and activates only a few experts (here: 2)

Fine-grained MoE uses more but smaller experts, and activates more experts (here: 4)

MoE with shared expert: also uses many small experts, but adds a shared expert that is always active




Figure 6: An annotated figure from "DeepSeekMoE: Towards Ultimate Expert Specialization in Mixture-of-Experts Language Models", <https://arxiv.org/abs/2401.06066>

The benefit of having a shared expert was first noted in the [DeepSpeedMoE paper](#), where they found that it boosts overall modeling performance compared to no shared experts. This is likely because common or repeated patterns don't have to be learned by multiple individual experts, which leaves them with more room for learning more specialized patterns.

1.3 DeepSeek Summary

To summarize, DeepSeek V3 is a massive 671-billion-parameter model that, at launch, outperformed other open-weight models, including the 405B Llama 3. Despite being larger, it is much more efficient at inference time thanks to its Mixture-of-Experts (MoE) architecture, which activates only a small subset of (just 37B) parameters per token.

Another key distinguishing feature is DeepSeek V3's use of Multi-Head Latent Attention (MLA) instead of Grouped-Query Attention (GQA). Both MLA and GQA are inference-efficient alternatives to standard Multi-Head Attention (MHA), particularly when using caching. While MLA is more complex to implement, a study in the DeepSeek-V2 paper shown it delivers better modeling performance than GQA.

2. OLMo 2

The OLMo series of models by the non-profit Allen Institute for AI is noteworthy due to transparency in terms of training data and code, as well as the relatively detailed technical reports.

While you probably won't find OLMo models at the top of any benchmark or leaderboard, they are pretty clean and, more importantly, a great blueprint for developing LLMs, thanks to their transparency.

And while OLMo models are popular because of their transparency, they are not that bad either. In fact, at the time of release in January (before Llama 4, Gemma 3, and Qwen 3), OLMo 2 models were sitting at the Pareto frontier of compute to performance, as shown in Figure 7 below.




Figure 7: Modeling benchmark performance (higher is better) vs pre-training cost (FLOPs; lower is better) for different LLMs. This is an annotated figure from the OLMo 2 paper, <https://arxiv.org/abs/2501.00656>

As mentioned earlier in this article, I aim to focus only on the LLM architecture details (not training or data) to keep it at a manageable length. So, what were the interesting architecture design choices in OLMo2? It mainly comes down to normalizations: the placement of RMSNorm layers as well as the addition of a QK-norm, which I will discuss below.

Another thing worth mentioning is that OLMo 2 still uses traditional Multi-Head Attention (MHA) instead of MLA or GQA.

2.1 Normalization Layer Placement

Overall, OLMo 2 largely follows the architecture of the original GPT model, similar to other contemporary LLMs. However, there are some noteworthy deviations. Let's start with the normalization layers.

Similar to Llama, Gemma, and most other LLMs, OLMo 2 switched from LayerNorm to RMSNorm.

But since RMSNorm is old hat (it's basically a simplified version of LayerNorm with fewer trainable parameters), I will skip the discussion of RMSNorm vs LayerNorm. (Curious readers can find an RMSNorm code implementation in my [GPT-2 to Llama conversion guide](#).)

However, it's worth discussing the placement of the RMSNorm layer. The original transformer (from the "[Attention is all you need](#)" paper) placed the two normalization layers in the transformer block *after* the attention module and the FeedForward module, respectively.

This is also known as Post-LN or Post-Norm.

GPT and most other LLMs that came after placed the normalization layers *before* the attention and FeedForward modules, which is known as Pre-LN or Pre-Norm. A comparison between Post- and Pre-Norm is shown in the figure below.




Figure 8: A comparison of *Post-Norm*, *Pre-Norm*, and OLMo 2's flavor of *Post-Norm*.

In [2020, Xiong et al.](#) showed that Pre-LN results in more well-behaved gradients at initialization. Furthermore, the researchers mentioned that Pre-LN even works well with careful learning rate warm-up, which is otherwise a crucial tool for Post-LN.

Now, the reason I am mentioning that is that OLMo 2 adopted a form of Post-LN (but with RMSNorm instead of LayerNorm, so I am calling it *Post-Norm*).

In OLMo 2, instead of placing the normalization layers before the attention and FeedFw layers, they place them after, as shown in the figure above. However, notice that in contrast to the original transformer architecture, the normalization layers are still inside the residual layers (skip connections).

So, why did they move the position of the normalization layers? The reason is that it helps with training stability, as shown in the figure below.




Figure 9: A plot showing the training stability for Pre-Norm (like in GPT-2, Llama 3, and many others) versus OLMo 2's flavor of Post-Norm. This is an annotated figure from the OLMo 2 paper, <https://arxiv.org/abs/2501.00656>

Unfortunately this figure shows the results of the reordering together with QK-Norm, w a separate concept. So, it's hard to tell how much the normalization layer reordering contributed by itself.

2.2 QK-Norm

Since the previous section already mentioned the QK-norm, and other LLMs we discuss such as Gemma 2 and Gemma 3, also use QK-norm, let's briefly discuss what this is.

QK-Norm is essentially yet another RMSNorm layer. It's placed inside the Multi-Head Attention (MHA) module and applied to the queries (q) and keys (k) before applying RoPE. To illustrate this, below is an excerpt of a Grouped-Query Attention (GQA) layer I wrote for [Qwen3 from-scratch implementation](#) (the QK-norm application in GQA is similar to MHA in OLMo):

```
class GroupedQueryAttention(nn.Module):
    def __init__(
        self, d_in, num_heads, num_kv_groups,
        head_dim=None, qk_norm=False, dtype=None
    ):
        # ...

        if qk_norm:
```

```

        self.q_norm = RMSNorm(head_dim, eps=1e-6)
        self.k_norm = RMSNorm(head_dim, eps=1e-6)
    else:
        self.q_norm = self.k_norm = None

def forward(self, x, mask, cos, sin):
    b, num_tokens, _ = x.shape

    # Apply projections
    queries = self.W_query(x)
    keys = self.W_key(x)
    values = self.W_value(x)

    # ...

    # Optional normalization
    if self.q_norm:
        queries = self.q_norm(queries)
    if self.k_norm:
        keys = self.k_norm(keys)

    # Apply RoPE
    queries = apply_rope(queries, cos, sin)
    keys = apply_rope(keys, cos, sin)

    # Expand K and V to match number of heads
    keys = keys.repeat_interleave(self.group_size, dim=1)
    values = values.repeat_interleave(self.group_size, dim=1)

    # Attention
    attn_scores = queries @ keys.transpose(2, 3)
    # ...

```

As mentioned earlier, together with Post-Norm, QK-Norm stabilizes the training. Note that QK-Norm was not invented by OLMo 2 but goes back to the [2023 Scaling Vision Transformers paper](#).

2.3 OLMo 2 Summary

In short, the noteworthy OLMo 2 architecture design decisions are primarily the RMSNorm placements: RMSNorm after instead of before the attention and FeedForward modules (flavor of Post-Norm), as well as the addition of RMSNorm for the queries and keys inside the attention mechanism (QK-Norm), which both, together, help stabilize the training loss.

Below is a figure that further compares OLMo 2 to Llama 3 side by side; as one can see, the architectures are otherwise relatively similar except for the fact that OLMo 2 still uses the traditional MHA instead of GQA. (However, the [OLMo 2 team released a 32B variant](#) 3 months later that uses GQA.)




Figure 10: An architecture comparison between Llama 3 and OLMo 2.

3. Gemma 3

Google's Gemma models have always been really good, and I think they have always been a bit underhyped compared to other popular models, like the Llama series.

One of the distinguishing aspects of Gemma is the rather large vocabulary size (to support multiple languages better), and the stronger focus on the 27B size (versus 8B or 70B).

note that Gemma 2 also comes in smaller sizes: 1B, 4B, and 12B.

The 27B size hits a really nice sweet spot: it's much more capable than an 8B model but as resource-intensive as a 70B model, and it runs just fine locally on my Mac Mini.

So, what else is interesting in [Gemma 3](#)? As discussed earlier, other models like DeepSeek V3/R1 use a Mixture-of-Experts (MoE) architecture to reduce memory requirements at inference, given a fixed model size. (The MoE approach is also used by several other models we will discuss later.)

Gemma 3 uses a different "trick" to reduce computational costs, namely sliding window attention.

3.1 Sliding Window Attention

With sliding window attention (originally introduced in the [LongFormer paper in 2020](#), and also already used by [Gemma 2](#)), the Gemma 3 team was able to reduce the memory requirements in the KV cache by a substantial amount, as shown in the figure below.




Figure 11: An annotated figure from Gemma 3 paper (<https://arxiv.org/abs/2503.19786>) showing the KV cache memory savings via sliding window attention.

So, what is sliding window attention? If we think of regular self-attention as a *global* attention mechanism, since each sequence element can access every other sequence element, we can think of sliding window attention as *local* attention, because here we restrict the context size around the current query position. This is illustrated in the figure below.




Figure 12: A comparison between regular attention (left) and sliding window attention (right).

Please note that sliding window attention can be used with both Multi-Head Attention Grouped-Query Attention; Gemma 3 uses grouped-query attention.

As mentioned above, sliding window attention is also referred to as *local* attention because the local window surrounds and moves with the current query position. In contrast, regular attention is *global* as each token can access all other tokens.

Now, as briefly mentioned above, the Gemma 2 predecessor architecture also used sliding window attention before. The difference in Gemma 3 is that they adjusted the ratio between global (regular) and local (sliding) attention.

For instance, Gemma 2 uses a hybrid attention mechanism that combines sliding window (local) and global attention in a 1:1 ratio. Each token can attend to a 4k-token window of nearby context.

Where Gemma 2 used sliding window attention in every other layer, Gemma 3 now has a 1:5 ratio, meaning there's only 1 full attention layer for every 5 sliding windows (local) attention layers; moreover, the sliding window size was reduced from 4096 (Gemma 2) to just 1024 (Gemma 3). This shifts the model's focus towards more efficient, localized computation.

According to their ablation study, the use of sliding window attention has minimal impact on modeling performance, as shown in the figure below.




Figure 13: An annotated figure from Gemma 3 paper (<https://arxiv.org/abs/2503.19786>) showing that sliding window attention has little to no impact on the LLM-generated output perplexity.

While sliding window attention is the most notable architecture aspect of Gemma 3, I will also briefly go over the placement of the normalization layers as a follow-up to the previous OLMo 2 section.

3.2 Normalization Layer Placement in Gemma

A small but interesting tidbit to highlight is that Gemma 3 uses RMSNorm in both a Pre and Post-Norm setting around its grouped-query attention module.

This is similar to Gemma 2 but still worth highlighting, as it differs from (1) the Post-Norm used in the original transformer ("Attention is all you need"), (2) the Pre-Norm, which was popularized by GPT-2 and used in many other architectures afterwards, and (3) the Post-Norm flavor in OLMo 2 that we saw earlier.




Figure 14: An architecture comparison between OLMo2 and Gemma 3; note the additional normalization layers in Gemma 3.

I think this normalization layer placement is a relatively intuitive approach as it gets the of both worlds: Pre-Norm and Post-Norm. In my opinion, a bit of extra normalization ca hurt. In the worst case, if the extra normalization is redundant, this adds a bit of inefficie through redundancy. In practice, since RMSNorm is relatively cheap in the grand scher things, this shouldn't have any noticeable impact, though.

3.3 Gemma 3 Summary

Gemma 3 is a well-performing open-weight LLM that, in my opinion, is a bit underappreciated in the open-source circles. The most interesting part is the use of sli

window attention to improve efficiency (it will be interesting to combine it with MoE in the future).

Also, Gemma 3 has a unique normalization layer placement, placing RMSNorm layers before and after the attention and FeedForward modules.

3.4 Bonus: Gemma 3n

A few months after the Gemma 3 release, Google shared [Gemma 3n](#), which is a Gemma model that has been optimized for small-device efficiency with the goal of running on phones.

One of the changes in Gemma 3n to achieve better efficiency is the so-called Per-Layer Embedding (PLE) parameters layer. The key idea here is to keep only a subset of the model parameters in GPU memory. Token-layer specific embeddings, such as those for text, audio, and vision modalities, are then streamed from the CPU or SSD on demand.

The figure below illustrates the PLE memory savings, listing 5.44 billion parameters for the standard Gemma 3 model. This likely refers to the Gemma 3 4-billion variant.




Figure 15: An annotated figure from Google's Gemma 3n blog (<https://developers.googleblog.com/en/introducing-gemma-3n/>) illustrating the PLE memory savings.

The 5.44 vs. 4 billion parameter discrepancy is because Google has an interesting way of reporting parameter counts in LLMs. They often exclude embedding parameters to make the model appear smaller, except in cases like this, where it is convenient to include them to make the model appear larger. This is not unique to Google, as this approach has become common practice across the field.

Another interesting trick is the [MatFormer](#) concept (short for Matryoshka Transformer). In this instance, Gemma 3n uses a single shared LLM (transformer) architecture that can be sliced into smaller, independently usable models. Each slice is trained to function on its own, and at inference time, we can run just the part you need (instead of the large model).

4. Mistral Small 3.1

[Mistral Small 3.1 24B](#), which was released in March shortly after Gemma 3, is noteworthy for outperforming Gemma 3 27B on several benchmarks (except for math) while being faster.

The reasons for the lower inference latency of Mistral Small 3.1 over Gemma 3 are likely due to their custom tokenizer, as well as shrinking the KV cache and layer count. Otherwise, it follows a standard architecture as shown in the figure below.




Figure 16: An architecture comparison between Gemma 3 27B and Mistral 3.1 Small 24B.

Interestingly, earlier Mistral models had utilized sliding window attention, but they appear to have abandoned it in Mistral Small 3.1 if we consider the default setting (`"sliding_window" : null`) in the official [Model Hub configuration file](#). Also, the [model card](#) makes no mention of it.

So, since Mistral uses regular Grouped-Query Attention instead of Grouped-Query Attention with a sliding window as in Gemma 3, maybe there are additional inference compute savings due to being able to use more optimized code (i.e., FlashAttention). For instance, I speculate that while sliding window attention reduces memory usage, it doesn't necessarily reduce inference latency, which is what Mistral Small 3.1 is focused on.

5. Llama 4

The extensive introductory discussion on Mixture-of-Experts (MoE) earlier in this article off again. [Llama 4](#) has also adopted an MoE approach and otherwise follows a relative standard architecture that is very similar to DeepSeek V3, as shown in the figure below (Llama 4 includes native multimodal support, similar to models like Gemma and Mistral). However, since this article focuses on language modeling, we only focus on the text model.




Figure 17: An architecture comparison between DeepSeek V3 (671-billion parameters) and Llama 4 Maverick (400-billion parameters).

While the Llama 4 Maverick architecture looks very similar to DeepSeek V3 overall, there are some interesting differences worth highlighting.

First, Llama 4 uses Grouped-Query Attention similar to its predecessors, whereas DeepSeek V3 uses Multi-Head Latent Attention, which we discussed at the beginning of this article. Now, both DeepSeek V3 and Llama 4 Maverick are very large architectures, with DeepSeek V3 being approximately 68% larger in its total parameter count. However, with 37 billion active parameters, DeepSeek V3 has more than twice as many active parameters as Llama 4 Maverick (17B).

Llama 4 Maverick uses a more classic MoE setup with fewer but larger experts (2 active experts with 8,192 hidden size each) compared to DeepSeek V3 (9 active experts with 2,048 hidden size each). Also, DeepSeek uses MoE layers in each transformer block (except the first 3), whereas Llama 4 alternates MoE and dense modules in every other transformer block.

Given the many small differences between architectures, it is difficult to determine their impact on final model performance. The main takeaway, however, is that MoE architectures have seen a significant rise in popularity in 2025.

6. Qwen3

The Qwen team consistently delivers high-quality open-weight LLMs. When I helped co-advise on the LLM efficiency challenge at NeurIPS 2023, I remember that the top winning solutions were all Qwen2-based.

Now, Qwen3 is another hit model series at the top of the leaderboards for their size class. There are 7 dense models: 0.6B, 1.7B, 4B, 8B, 14B, and 32B. And there are 2 MoE models: 30B-A3B, and 235B-A22B.

(By the way, note that the missing whitespace in "Qwen3" is not a typo; I simply try to preserve the original spelling the Qwen developers chose.)

6.1 Qwen3 (Dense)

Let's discuss the dense model architecture first. As of this writing, the 0.6B model may be the smallest current-generation open-weight model out there. And based on my personal experience, it performs really well given its small size. It has great token/sec throughput and low memory footprint if you are planning to run it locally. But what's more, it's also easy to train locally (for educational purposes) due to its small size.

So, Qwen3 0.6B has replaced Llama 3 1B for me for most purposes. A comparison between these two architectures is shown below.




Figure 18: An architecture comparison between Qwen3 0.6B and Llama 3 1B; notice that Qwen3 is a deeper architecture with more layers, whereas Llama 3 is a wider architecture with more attention heads.

If you are interested in a human-readable Qwen3 implementation without external third party LLM library dependencies, I recently implemented [Qwen3 from scratch \(in pure PyTorch\)](#).

The computational performance numbers in the figure above are based on my from-scratch PyTorch implementations when run on an A100 GPU. As one can see, Qwen3 has a smaller memory footprint as it is a smaller architecture overall, but also uses smaller hidden layer and fewer attention heads. However, it uses more transformer blocks than Llama 3, which leads to a slower runtime (lower tokens/sec generation speed).

6.2 Qwen3 (MoE)

As mentioned earlier, Qwen3 also comes in two MoE flavors: 30B-A3B and 235B-A22. Why do some architectures, like Qwen3, come as regular (dense) and MoE (sparse) variants?

As mentioned at the beginning of this article, MoE variants help reduce inference costs large base models. Offering both dense and MoE versions gives users flexibility depend on their goals and constraints.

Dense models are typically more straightforward to fine-tune, deploy, and optimize across various hardware.

On the other hand, MoE models are optimized for scaling inference. For instance, at a fixed inference budget, they can achieve a higher overall model capacity (i.e., knowledge up to during training due to being larger) without proportionally increasing inference costs.

By releasing both types, the Qwen3 series can support a broader range of use cases: dense models for robustness, simplicity, and fine-tuning, and MoE models for efficient serving at scale.

To round up this section, let's look at Qwen3 235B-A22B (note that the A22B stands for "22B active parameters") to DeepSeek V3, which has almost twice as many active parameters (37B).




Figure 19: An architecture comparison between DeepSeek V3 and Qwen3 235B-A22B.

As shown in the figure above, the DeepSeek V3 and Qwen3 235B-A22B architectures remarkably similar. What's noteworthy, though, is that the Qwen3 model moved away from using a shared expert (earlier Qwen models, such as [Qwen2.5-MoE](#) did use a shared expert).

Unfortunately, the Qwen3 team did not disclose any reason as to why they moved away from shared experts. If I had to guess, it was perhaps simply not necessary for training stability in their setup when they increased the experts from 2 (in Qwen2.5-MoE) to 8 (in Qwen3), then they were able to save the extra compute/memory cost by using only 8 instead of 2. (However, this doesn't explain why DeepSeek V3 is still keeping their shared expert.)

Update. [Junyang Lin](#), one of the developers of Qwen3, responded as follows:

At that moment we did not find significant enough improvement on shared expert architecture. We were worrying about the optimization for inference caused by shared expert. No straightforward answer to this question honestly.

7. SmolLM3

[SmolLM3](#) is perhaps not as nearly as popular as the other LLMs covered in this article, but I thought it is still an interesting model to include as it offers really good modeling performance at a relatively small and convenient 3-billion parameter model size that sits between the 1.1B and 4B Qwen3 model, as shown in the figure below.

Moreover, it also shared a lot of the training details, similar to OLMo, which is rare and appreciated!




Figure 20: An annotated figure from the SmoLLM3 announcement post, <https://huggingface.co/blog/smolLM3>, comparing the SmoLLM3 win rate to Qwen3 1.7B and 4B as well as Llama 3 3B and Gemma 3 4B.

As shown in the architecture comparison figure below, the SmoLLM3 architecture looks standard. The perhaps most interesting aspect is its use of NoPE (No Positional Embeddings), though.




Figure 21: A side-by-side architecture comparison between Qwen3 4B and SmolLM3 3B.

7.1 No Positional Embeddings (NoPE)

NoPE is, in LLM contexts, an older idea that goes back to a 2023 paper ([The Impact of Positional Encoding on Length Generalization in Transformers](#)) to remove explicit position information injection (like through classic absolute positional embedding layers in early architectures or nowadays RoPE).

In transformer-based LLMs, positional encoding is typically necessary because self-attention treats tokens independently of order. Absolute position embeddings solve this by adding an additional embedding layer that adds information to the token embeddings.




Figure 22: A modified figure from my Build A Large Language Model (From Scratch) book (<https://www.amazon.com/Build-Large-Language-Model-Scratch/dp/1633437167>) illustrating absolute positional embeddings.

RoPE, on the other hand, solves this by rotating the query and key vectors relative to the token position.

In NoPE layers, however, no such positional signal is added at all: not fixed, not learned, relative. Nothing.

Even though there is no positional embedding, the model still knows which tokens come before, thanks to the causal attention mask. This mask prevents each token from attending to future ones. As a result, a token at position t can only see tokens at positions $\leq t$, which preserves the autoregressive ordering.

So while there is no positional information that is explicitly added, there is still an implicit sense of direction baked into the model's structure, and the LLM, in the regular gradient descent-based training, can learn to exploit it if it finds it beneficial for the optimization objective. (Check out the NoPE paper's theorems for more information.)

So, overall, the [NoPE paper](#) not only found that no positional information injection is necessary, but it also found that NoPE has better length generalization, which means that LLM answering performance deteriorates less with increased sequence length, as shown in the figure below.




Figure 23: An annotated figure from the NoPE paper (<https://arxiv.org/abs/2305.19466>) showing better length generalization with NoPE.

Note that the experiments shown above were conducted with a relatively small GPT-style model of approximately 100 million parameters and relatively small context sizes. It is unclear how well these findings generalize to larger, contemporary LLMs.

For this reason, the SmolLM3 team likely only "applied" NoPE (or rather omitted RoPE) every 4th layer.

8. Kimi K2 and Kimi K2 Thinking

[Kimi K2](#) recently made big waves in the AI community due to being an open-weight model with an incredibly good performance. According to benchmarks, it's on par with the best proprietary models like Google's Gemini, Anthropic's Claude, and OpenAI's ChatGPT models.


A notable aspect is its use of a variant of the relatively new [Muon](#) optimizer over AdamW. As far as I know, this is the first time Muon was used over AdamW for any production model.

this size ([previously](#), it has only been shown to scale up to 16B). This resulted in very nice training loss curves, which probably helped catapult this model to the top of the aforementioned benchmarks.

While people commented that the loss was exceptionally smooth (due to the lack of spikes), I think it's not exceptionally smooth (e.g., see the OLMo 2 loss curve in the figure below). The L2 norm of the gradient would probably be a better metric to track training stability. However, what's remarkable is how well the loss curve decays.

However, as mentioned in the introduction of this article, training methodologies are a topic for another time.

Kimi K2 training loss curve



OLMo & OLMo 2 training loss curves




Figure 24: Annotated figures from the Kimi K2 announcement blog article (<https://moonshotai.github.io/Kimi-K2/>) and the OLMo 2 paper (<https://arxiv.org/abs/2305.19466>).

It may be the biggest LLM of this generation as of this writing (given the constraints that Llama 4 Behemoth is not released, proprietary LLMs don't count, and Google's 1.6 trillion [Switch Transformer](#) is an encoder-decoder architecture from a different generation).

DeepSeek V3/R1 (671 billion)
more heads, fewer experts

Kimi K2 (1 trillion)
fewer heads, more experts

DeepSeek V3/R1 Architecture Details:

- Vocabulary size of 129k
- Supported context length of 128k tokens
- First 3 blocks use dense FFN with hidden size 18,432 instead of MoE
- Sample input text → Tokenized text → Token embedding layer (Embedding dimension of 7,168)
- 61x (61 blocks)
- 128 heads
- Resource savings:
 - Model size is 671B
 - but only 1 (shared) + 8 experts active per token
 - only 37B parameters are active per inference step

Kimi K2 Architecture Details:

- Vocabulary size of 160k
- Supported context length of 128k tokens
- First 1 block uses dense FFN with hidden size 18,432 instead of MoE
- Sample input text → Tokenized text → Token embedding layer (Embedding dimension of 7,168)
- 61x (61 blocks)
- 64 heads
- Resource saving:
 - Model size is 1T
 - but only 1 (shared) + 8 experts active per token
 - only 32B parameters are active per inference step

As shown in the figure above, Kimi K2 is basically the same as DeepSeek V3, except it uses more experts in the MoE modules and fewer heads in the Multi-head Latent Attention (MLA) module.

So, most likely the Kimi K2 team took these lessons to heart and shared Kimi K2 as an open weight model, before DeepSeek R2 was released. As of this writing, Kimi K2 is the most

impressive open-weight model.

Update: On Nov 6, 2025 the Kimi K2 team also released their new "Thinking" model v. The architecture is unchanged from Kimi K2 above, except that they extended the context size from 128k to 256k.

According to the [benchmarks shared by the Kimi team](#), the model exceeds the performance of the leading proprietary LLMs. (Unfortunately, there is no direct comparison to DeepSeek R1.





Figure 25.2: DeepSeek R1 versus Kimi K2 Thinking architecture (top) and Kimi K2 Thinking benchmarks (bottom).

9. GPT-OSS

OpenAI's [released](#) gpt-oss-120b and gpt-oss-20b, their first open-weight models since GPT-2 in 2019, about one week after I wrote this article. Since OpenAI's open-weight models have been so widely anticipated, I updated this article to include them. I will keep this s

brief, but I have written another, much more detailed article dedicated to the gpt-oss m here:



From GPT-2 to gpt-oss: Analyzing the Architectural Advances

SEBASTIAN RASCHKA, PHD · 2025年8月9日

[Read full story →](#)

Before summarizing the interesting tidbits, let's start with an overview of the two mode gpt-oss-20b and gpt-oss-120b, as shown in Figure 26 below.




Figure 26: Architecture overview of the two gpt-oss models.

Looking at Figure 26, the architecture contains all the familiar components we have seen in other architectures discussed previously. For instance, Figure 27 puts the smaller gpt-oss-20b architecture next to Qwen3 30B-A3B, which is also an MoE model with a similar number of active parameters (gpt-oss has 3.6B active parameters, and Qwen3 30B-A3B has 3.3B active parameters).




Figure 27: Architecture comparison between gpt-oss and Qwen3

One aspect not shown in Figure 27 is that gpt-oss uses sliding window attention (similar to Gemma 3, but in every other layer instead of using a 5:1 ratio).

9.1 Width Versus Depth

Figure 27 shows that gpt-oss and Qwen3 use similar components. But if we look at the models closely, we see that Qwen3 is a much deeper architecture with its 48 transformer blocks instead of 24.

On the other hand, gpt-oss is a much wider architecture:

- An embedding dimension of 2880 instead of 2048
- An intermediate expert (feed forward) projection dimension of also 2880 instead of 2048

It's also worth noting that gpt-oss uses twice as many attention heads, but this doesn't directly increase the model's width. The width is determined by the embedding dimension.

Does one approach offer advantages over the other given a fixed number of parameters? As a rule of thumb, deeper models have more flexibility but can be harder to train due to

instability issues, due to exploding and vanishing gradients (which RMSNorm and short connections aim to mitigate).

Wider architectures have the advantage of being faster during inference (with a higher tokens/second throughput) due to better parallelization at a higher memory cost.

When it comes to modeling performance, there's unfortunately no good apples-to-apples comparison I am aware of (where parameter size and datasets are kept constant) except an ablation study in the [Gemma 2 paper \(Table 9\)](#), which found that for a 9B parameter architecture, a wider setup is slightly better than a deeper setup. Across 4 benchmarks wider model achieved a 52.0 average score, and the deeper model achieved a 50.8 average score.

9.2 Few Large Versus Many Small Experts

As shown in Figure 27 above, it's also noteworthy that gpt-oss has a surprisingly small number of experts (32 instead of 128), and only uses 4 instead of 8 active experts per token. However, each expert is much larger than the experts in Qwen3.

This is interesting because the recent trends and developments point towards more, smaller models as being beneficial. This change, at a constant total parameter size, is nicely illustrated in Figure 28 below from the DeepSeekMoE paper.

Early MoE: Has bigger and fewer experts, and activates only a few experts (here: 2)

Fine-grained MoE uses more but smaller experts, and activates more experts (here: 4)

MoE with shared expert: also uses many small experts, but adds a shared expert that is always active




Figure 28: An annotated figure from "DeepSeekMoE: Towards Ultimate Expert Specialization in Mixture-of-Experts Language Models", <https://arxiv.org/abs/2401.06066>

Notably, unlike DeepSeek's models, neither gpt-oss nor Qwen3 uses shared experts, though.

9.3 Attention Bias and Attention Sinks

Both gpt-oss and Qwen3 use grouped query attention. The main difference is that gpt restricts the context size via sliding window attention in each second layer, as mentioned earlier.

However, there's one interesting detail that caught my eye. It seems that gpt-oss uses units for the attention weights, as shown in Figure 29 below.



Figure 29: gpt-oss models use bias units in the attention layers. See code example here.

I haven't seen these bias units being used since the GPT-2 days, and they are common regarded as redundant. Indeed, I found a recent paper that shows mathematically that at least true for the key transformation (k_proj). Furthermore, the empirical results show there is little difference between with and without bias units (see Figure 30 below).

	GPT-2	RoBERTa-base
Original	2.9251	5.8890
No b_k	2.9250	5.8909

Test set loss, lower is better

Figure 30: Table from <https://arxiv.org/pdf/2302.08626> showing the average test loss when the models were trained from scratch with and without bias units.

Another detail you may have noticed is the definition of `sinks` in the code screenshot in Figure 30. In general models, attention sinks are special "always-attended" tokens placed at the start of the sequence to stabilize attention, which is especially useful in long-context scenarios. I.e., if the context gets very long, this special attended token at the beginning is always attended to, and it can learn to store some generally useful information about the entire sequence. (I think it was originally proposed in the [Efficient Streaming Language Models: Attention Sinks](#) paper.)

In the gpt-oss implementation, *attention sinks* are not actual tokens in the input sequence. Instead, they are learned per-head bias logits that are appended to the attention score:

(Figure 31). The goal is the same as with the above-mentioned attention sinks, but with modifying the tokenized inputs.


```

228 class GptOssAttention(Qwen2Attention):
229     def __init__(self, config: GptOssConfig, layer_idx: int):
230         super().__init__(config, layer_idx)
231         self.q_proj = nn.Linear(
232             config.hidden_size, config.num_attention_heads * self.head_dim, bias=config.attention
233         )
234         self.k_proj = nn.Linear(
235             config.hidden_size, config.num_key_value_heads * self.head_dim, bias=config.attention
236         )
237         self.v_proj = nn.Linear(
238             config.hidden_size, config.num_key_value_heads * self.head_dim, bias=config.attention
239         )
240         self.o_proj = nn.Linear(
241             config.num_attention_heads * self.head_dim, config.hidden_size, bias=config.attention
242         )
243         self.sinks = nn.Parameter(torch.empty(config.num_attention_heads))
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000

```

Figure 31: The use of attention sinks in gpt-oss; based on the Hugging Face code [here](#).

For more information about gpt-oss, and how it compares to GPT-2, please see my other gpt-oss article:



From GPT-2 to gpt-oss: Analyzing the Architectural Advances

SEBASTIAN RASCHKA, PHD • 2025年8月9日

[Read full story →](#)

10. Grok 2.5

A few weeks after this article first went online, xAI released the weights of their 270B-parameter Grok 2.5 model.

I thought it would be worth including here, since Grok 2.5 was xAI's flagship production model last year. Up to this point, all models we discussed were released as open-weight models from the start. For example, gpt-oss is likely not an open-weight clone of GPT- rather a custom model trained specifically for the open-source community.

With Grok 2.5, we get a rare look at a real production system, even if it is last year's.

Architecturally, Grok 2.5 looks fairly standard overall (Figure 32), but there are a few noteworthy details.




Figure 32: Grok 2.5 next to a Qwen3 model of comparable size

For instance, Grok 2.5 uses a small number of large experts (eight), which reflects an old trend. As discussed earlier, more recent designs such as those in the DeepSeekMoE paper favor a larger number of smaller experts (this is also present in Qwen3).

Another interesting choice is the use of what amounts to a shared expert. The addition of the SwiGLU module shown on the left in Figure 32 functions as an always-on, shared expert. This is not identical to the classic shared-expert design since its intermediate dimension is different.

but the idea is the same. (I still find it interesting that Qwen3 omitted shared experts, as will be interesting to see if that changes with Qwen4 and later models.)


11. GLM-4.5

[GLM-4.5](#) is another major release this year.

It is an instruction/reasoning hybrid similar to Qwen3, but even better optimized for fun calling and agent-style contexts.

LLM Performance Evaluation: Agentic, Reasoning, And Coding Benchmarks

12 benchmarks: MMLU-Pro, AIME 24, MATH-500, SciCode, GPQA, HLE, LCB (2407-2501), SWE-Bench Verified, Terminal-Bench, TAU-Bench, BFCL V3, Brow




Agentic

Agentic Benchmarks: TAU-Bench, BFCL V3 (Full), BrowseComp



Reasoning

Reasoning Benchmarks: MMLU-Pro, AIME 24, MATH 500, SciCode, GPQA, HLE, LCB (2407-2501)



Coding

Coding Benchmarks: SWE-Bench Verified, Terminal-Bench




Figure 33: GLM-4.5 benchmark from the official GitHub repository at <https://github.com/zai-org/GLM-4.5>

As shown in Figure 34, GLM-4.5 comes in two variants. The flagship 355-billion-parameter model outperforms Claude 4 Opus on average across 12 benchmarks and trails only slightly behind OpenAI's o3 and xAI's Grok 4. There is also GLM-4.5-Air, a more compact 106

billion-parameter version that delivers performance only marginally below the 355-billion model.

Figure 35 compares the 355-billion architecture to Qwen3.




Figure 34: GLM-4.5 next to a similarly-sized Qwen3 model.

The designs are largely similar, but GLM-4.5 adopts a structural choice first introduced by DeepSeek V3: 3 dense layers precede the Mixture-of-Experts (MoE) blocks. Why? Starting with several dense layers improves convergence stability and overall performance in large MoE systems. If MoE routing is introduced immediately, the instability of sparse expert selection can interfere with early syntactic and semantic feature extraction. So, one might argue that by keeping the initial layers dense ensures the model forms stable low-level representations before routing decisions begin to shape higher-level processing.

Also, GLM-4.5 uses a shared expert similar to DeepSeek V3 (and unlike Qwen3).

(Interestingly, GLM-4.5 also retains the attention bias mechanism used in GPT-2 and GPT-3.)


12. Qwen3-Next

On 11 September 2025, the Qwen3 team released Qwen3 Next 80B-A3B (Figure 35), available in both Instruct and Thinking variants. While its design builds on the previous discussed Qwen3 architecture, I included it here as a separate entry to keep the figure numbering consistent and to draw attention to some of its design changes.

12.1 Expert Size and Number

The new Qwen3 Next architecture stands out because, despite being 3× smaller than previous 235B-A22B model (Figure 35), it introduces four times as many experts and adds a shared expert. Both of these design choices (a high expert count and the inclusion of a shared expert) were future directions I had highlighted prior to this release, particularly in the video version of the article that I linked at the top.

Qwen3 235B-A22B



Qwen3 Next 80B-A3B




Figure 35: The original Qwen3 model released in May (left) next to the Qwen3 Next model released in September (right).

12.2 Gated DeltaNet + Gated Attention Hybrid

The other highlight is that they replace the regular attention mechanism by a [Gated Delta](#) + [Gated Attention](#) hybrid, which helps enable the native 262k token context length in terms of memory usage (the previous 235B-A22B model model supported 32k natively, and with [YaRN](#) scaling.)

So how does this new attention hybrid work? Compared to grouped-query attention (which is still standard scaled dot-product attention (sharing K/V across query-head groups to cut KV-cache size and memory bandwidth as discussed earlier but whose decode cost and cache still grow with sequence length)), their hybrid mechanism mixes *Gated Delta* blocks with *Gated Attention* blocks with in a 3:1 ratio as shown in Figure 36.




Figure 36: The Gated DeltaNet + Gated Attention hybrid mechanism. Note that these are arranged in a 3:1 ratio, meaning that 3 transformer blocks with Gated DeltaNet are followed by 1 transformer block with Gated Attention. The right subfigure is from the official Qwen3 blog: <https://qwen.ai/blog?id=4074cca80393150c248e508aa62983f9cb7d27cd&from=research.latest-advancements-list>

We can think of the gated attention block as standard scaled-dot-product attention that can be used in GQA, but it has a few tweaks on top. The main differences between *gated attention* and plain GQA block are:

1. an output gate (sigmoid-controlled, usually per-channel) that scales the attention before it is added back to the residual;
2. zero-centered RMSNorm for QKNorm, rather than a standard RMSNorm;
3. partial RoPE (on a subset of dimensions).

Note that these are essentially just stability changes to GQA.

The Gated DeltaNet is a more significant change. In the DeltaNet block, q , k , v and two (α, β) are produced by linear and lightweight convolutional layers with normalization, and the attention layer replaces attention with a fast-weight [delta rule](#) update.

However, the tradeoff is that DeltaNet offers less precise content-based retrieval than attention, which is why one gated attention layer remains.

Given that attention grows quadratically, the DeltaNet component was added to help with memory efficiency. In the "linear-time, cache-free" family, the DeltaNet block is a lesser alternative to Mamba. Mamba keeps a state with a learned state-space filter (essentially a dynamic convolution over time). DeltaNet keeps a tiny fast-weight memory updated with α and β and reads it with q , with small convolutions only used only to help form q , k , v , α , β .

12.3 Multi-Token Prediction

The two subsections above describe two design decisions geared towards efficiency. Since all good things come in threes, the Qwen3 also adds another efficiency-technique on top: [Multi-Token Prediction](#) (MTP).

(Note that DeepSeek V3 & V3.2, and later GLM-4.5 and MiniMax-M2 all use MTP during training; however, since it's a training technique, I haven't explicitly discussed it in the architecture comparisons.)

Multi-token prediction trains the LLM to predict several future tokens, instead of a single token at each step. Here, at each position t , small extra heads (linear layers) output logits for $t+1 \dots t+k$, and we sum cross-entropy losses for these offsets (in the [MTP](#) paper the researchers recommended $k=4$). This additional signal speeds up training, and inference may remain at generating one token at a time. However, the extra heads can be used in speculative multi-token decoding, which is what Qwen3-Next seems to do, however, the details are still a bit sparse:

Qwen3-Next introduces a native Multi-Token Prediction (MTP) mechanism, which not only yields an MTP module with a high acceptance rate for Speculative Decoding but also enhances the overall performance. Additionally, Qwen3-Next specifically optimizes multi-step inference performance of MTP, further improving the acceptance rate of Speculative Decoding in real scenarios through multi-step training that maintains consistency between training and inference. [Source: Qwen3-Next blog post](#)

12.4 Qwen3-Coder-Next

In early February 2026, the Qwen3 team [shared](#) the 80B Qwen3-Coder-Next model (80B parameters active), which made big headlines for outperforming much larger models like DeepSeek V3.2 (37B active) and Kimi K2.5 and GLM-7.5 (both 32B active) on coding

Moreover, the Qwen3-Coder-Next SWE-Bench Pro Performance is roughly on par with Claude-Sonnet-4.5 (and only slightly below Claude-Opus-4.5), which is impressive for an open-weight model!

Note that the architecture behind Qwen3-Coder-Next is exactly the same as Qwen3-80B, which we discussed above (in fact, they used Qwen3-Next as a base model to train Qwen3-Coder-Next. Since this is an article about LLM architectures, the training details are outside the scope. However, interested readers can find more information in their [detailed technical report on GitHub](#).

13. MiniMax-M2

Recently, open-weight LLM developers shared flavors of their core architectures optimized for efficiency. One example is Qwen3-Next (see previous section), which replaces some of the full attention blocks with a fast gated DeltaNet module. Another example is DeepSeek V3.2, which uses sparse attention, a linear attention variant that trades off some model performance for improved computational performance (I plan to cover this mechanism in more detail in an upcoming article).

Now, [MiniMax-M1](#) falls into a similar category to the models above, in that it uses a linear attention variant (lightning attention) that offers improved efficiency over regular (full) attention. I originally didn't cover MiniMax M1 as it wasn't quite as popular as some of the other models discussed here. However, their new [MiniMax-M2](#) release is currently considered the best open-weight model (according to benchmark performance), which makes it too big to ignore.




Figure 37: MiniMax-M2 benchmark performance compared to other popular open-weight and proprietary LLMs. Image from the official model hub release [readme](#) file.

As shown in the overview figure below, I grouped MiniMax-M2 with the other decoder-transformer LLMs as it does not use the efficient lightning attention variant proposed in MiniMax-M1. Instead, the developers went back to using full attention, likely to improve modeling (and benchmark) performance.




Figure 38: A timeline of the main LLMs covered in this article, next to some of the attention-hybrid models that constitute more efficient alternatives, trading off some modeling performance with improved efficiency.

Overall, MiniMax-M2 is surprisingly similar to Qwen3. Besides changing the number of layers, sizes, etc., it uses the same components overall.

13.1 Per-Layer QK-Norm

Perhaps the one noteworthy highlight here is that MiniMax-M2 uses a so-called “per_layer QK-Norm instead of the regular QK-Norm. A closer look at the [code](#) reveals that it is implemented like this inside the attention mechanism:

```
self.q_norm = MiniMaxText01RMSNormTP(self.head_dim * self.total_num_heads, eps=...)

self.k_norm = MiniMaxText01RMSNormTP(self.head_dim * self.total_num_kv_heads, eps=...)
```


Here, the `hidden_size` equals the concatenated heads (`num_heads * head_dim`), so that RMSNorm has a scale vector with distinct parameters for every head (and each head c

So, the “per_layer” means that the RMSNorm (used for QK-Norm as explained earlier defined in each transformer block (as in regular QK-Norm), but, in addition, instead of r it across attention heads, it’s a unique QK-Norm for each attention head.

The [model configuration file](#) also includes a sliding-window attention setting (similar to Gemma 3 in section 3), but, as in Mistral 3.1 (discussed in section 4), it is disabled by default.

Otherwise, besides the per-layer QK-Norm, the architecture is very similar to Qwen3, as shown in the figure below.

Qwen3 235B-A22B



MiniMax-M2 (230B)




Figure 39: Comparison between Qwen3 and MiniMax-M2.

13.2 MoE Sparsity

Other interesting tidbits, as shown in the figure below, include the fact that they don't use shared experts (similar to Qwen3 but unlike Qwen3-Next). As mentioned earlier, in my opinion, shared experts are useful because they reduce redundancy among the other experts.

Also, as apparent from the figure above, MiniMax-M2 is twice as "sparse" as Qwen3. I.e. roughly the same size as Qwen3 235B-A22B, MiniMax-M2 has only 10B instead of 22B active experts per token (that is, 4.37% of the parameters are used in each inference step in MiniMax-M2, whereas Qwen3 uses 9.36% active tokens).

13.3 Partial RoPE

Lastly, similar to MiniMax-M1, MiniMax-M2 uses a “partial” instead of regular RoPE inside attention modules to encode positional information. Similar to regular RoPE, the rotation is applied to the queries and keys after applying QK-Norm.

Partial RoPE here means only the first `rotary_dim` channels of each head get rotary position encodings, and the remaining `head_dim - rotary_dim` channels remain unchanged.

In the official M1 [README](#) file, the developers mention

Rotary Position Embedding (RoPE) applied to half of the attention head dimension with a base frequency of 10,000,000

We can picture it as follows:

```
Full RoPE:    [r r r r r r r r]
Partial RoPE: [r r r r — — — —]
```

where in the conceptual illustration above, the “r”s show rotated (position-encoded) dimensions, and the dashes are the untouched dimensions.

What’s the point of this? In the [M1 paper](#), the developers stated that

...implementing RoPE on half of the softmax attention dimensions enables length extrapolation without performance degradation.

My speculation is that this prevents “too much” rotation for long sequences, and particularly those that are longer than the longest documents in the training dataset. I.e., the rationale here could be that no rotation is better than a “bad” or “too extreme” rotation that the model hasn’t seen before in training.

14. Kimi Linear

There's recently been a revival in linear attention mechanisms to improve the efficiency LLMs.

The attention mechanism introduced in the Attention Is All You Need paper (2017), aka scaled-dot-product attention, remains the most popular attention variant in today's LLMs. Besides traditional multi-head attention, it's also used in the more efficient flavors like grouped-query attention, sliding window attention, and multi-head latent attention.


14.1 Traditional Attention and Quadratic Cost

The original attention mechanism scales quadratically with the sequence length:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right) V$$

This is because the query (Q), key (K), and value (V) are n -by- d matrices, where d is the embedding dimension (a hyperparameter) and n is the sequence length (i.e., the number of tokens).


You can find more details on attention in my other article:



Understanding and Coding Self-Attention, Multi-Head Attention, Causal-Attention, and Cross-Attention in LLMs

2024年1月14日

[Read full story →](#)

Figure 40: Illustration of the quadratic cost in attention due to sequence length n .

14.2 Linear attention

Linear attention variants have been around for a long time, and I remember seeing tons papers in the 2020s. For example, one of the earliest I recall is the 2020 [Transformers: RNNs: Fast Autoregressive Transformers with Linear Attention](#) paper, where the research approximated the attention mechanism:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V \approx \phi(Q)(\phi(K)^T V)$$

Here, $\phi(\cdot)$ is a kernel feature function, set to $\phi(x) = \text{elu}(x) + 1$.

This approximation is efficient because it avoids explicitly computing the $n \times n$ attention matrix QK^T . Instead of performing all pairwise token interactions (which cost $O(n^2)$ time and memory).

I don't want to dwell too long on these older attempts. But the bottom line was that they reduced both time and memory complexity from $O(n^2)$ to $O(n)$ to making attention more efficient for long sequences.

However, they never really gained traction as they degraded the model accuracy, and I never really seen one of these variants applied in an open-weight state-of-the-art LLM.

14.3 Linear Attention Revival

In the second half of this year, there was a bit of a revival of linear attention variants. The notable model was [MiniMax-M1](#) with lightning attention, a 456B parameter mixture-of-experts (MoE) model with 46B active parameters, which came out back in June.

Then, in August, the Qwen3 team followed up with Qwen3-Next, which I discussed in detail above. Then, in September, the DeepSeek Team announced DeepSeek V3.2. All models (MiniMax-M1, Qwen3-Next, DeepSeek V3.2) replace the traditional quadratic attention variants in most or all of their layers with efficient linear variants.

Interestingly, there was a recent plot twist, where the MiniMax team released their new parameter M2 model (discussed in section 13) without linear attention, going back to regular attention. The team stated that linear attention is tricky in production LLMs. It seemed to work fine with regular prompts, but it had poor accuracy in reasoning and multi-turn tasks which are not only important for regular chat sessions but also agentic applications.

This could have been a turning point where linear attention may not be worth pursuing all. However, it gets more interesting. In October, the Kimi team released their new [Kimi](#)

model with linear attention.





Figure 41: An overview of the linear attention hybrid architectures.

Side note: I could have grouped Qwen3-Next and Kimi Linear with the other transformer-state space model (SSM) hybrids in the overview figure. Personally, I see these other transformer-SSM hybrids as SSMs with transformer components, whereas I see the models discussed here (Qwen3-Next and Kimi Linear) as transformers with SSM components. However, since I have listed IBM Granite 4.0 and NVIDIA Nemotron Nano 2 in the transformer-SSM box, an argument could be made for putting them into a single category.

14.4 Kimi Linear vs. Qwen3-Next

Kimi Linear shares several structural similarities with Qwen3-Next. Both models rely on a hybrid attention strategy. Concretely, they combine lightweight linear attention with heavy full attention layers. Specifically, both use a 3:1 ratio, meaning for every three transformer blocks employing the linear Gated DeltaNet variant, there's one block that uses full attention as shown in the figure below.

Qwen3 Next 80B-A3B



Kimi Linear 48B-A3B




Figure 42: Qwen3-Next and Kimi Linear side by side.

Gated DeltaNet is a linear attention variant with inspiration from recurrent neural networks including a gating mechanism from the [Gated Delta Networks: Improving Mamba2 with Delta Rule](#) paper. In a sense, Gated DeltaNet is a DeltaNet with Mamba-style gating, and DeltaNet is a linear attention mechanism. Due to the overview-nature of this article, DeltaNet would be a good topic for a separate article in the future.

Note that the omission of the RoPE box in the Kimi Linear part of the figure above is intentional. Kimi applies NoPE (No Positional Embedding) in multi-head latent attention layers (global attention). As the authors state, this lets MLA run as pure multi-query attention at inference and avoids RoPE retuning for long-context scaling (the positional bias is supposedly handled by the Kimi Delta Attention blocks). For more information on MLA, multi-query attention, which is a special case of grouped-query attention, please see the [Big LLM Architecture Comparison](#) article.

In addition, I've written more about Gated DeltaNet [here](#).

14.5 Kimi Delta Attention

Kimi Linear modifies the linear attention mechanism of Qwen3-Next by the Kimi Delta Attention (KDA) mechanism, which is essentially a refinement of Gated DeltaNet. When Qwen3-Next applies a scalar gate (one value per attention head) to control the memory decay rate, Kimi Linear replaces it with a channel-wise gating for each feature dimension. According to the authors, this gives more control over the memory, and this, in turn, improves long-context reasoning.

In addition, for the full attention layers, Kimi Linear replaces Qwen3-Next's gated attention layers (which are essentially standard multi-head attention layers with output gating) with Multi-Head Latent Attention (MLA). This is the same MLA mechanism we discussed earlier in the DeepSeek V3/R1 section but with an additional gate. (To recap, MLA compresses the key/value space to reduce the KV cache size.)

There's no direct comparison to Qwen3-Next, but compared to the Gated DeltaNet-H model from the Gated DeltaNet paper (which is essentially Gated DeltaNet with sliding window attention), Kimi Linear achieves higher modeling accuracy while maintaining the same token-generation speed.




Figure 43: Annotated figure from the Kimi Linear paper showing that Kimi Linear is as fast as GatedDeltaNet, and much faster than an architecture with multi-head latent attention (like DeepSeek V3/R1), while having a higher benchmark performance.

Furthermore, according to the ablation studies in the [DeepSeek-V2 paper](#), MLA is on par with regular full attention when the hyperparameters are carefully chosen.

And the fact that Kimi Linear compares favorably to MLA on long-context and reasoning benchmarks makes linear attention variant once again promising for larger state-of-the-art models. That being said, Kimi Linear is 48B-parameter large, but it's 20x smaller than K2. It will be interesting to see if the Kimi team adopts this approach for their upcoming model.

15. Olmo 3 Thinking


Allen AI [released their new Olmo 3](#) 7B and 32B models on November 20. (The official spelling was changed from OLMo to Olmo, so I will be adopting that in this section.)

As mentioned earlier, Olmo models are always interesting because they are fully open-source. Here, that means that the team also shares [detailed training reports](#), multiple checkpoints, information about the training data, and so forth. In other words, Olmo models are fully transparent.

This time, the Olmo suite also comes in an additional reasoning model flavor (next to base and instruct models), and there are lots of interesting details about the training in Olmo [technical report](#). However, since this is an article about architectural comparisons, this section focuses only on Olmo 3's architecture.

The closest model to compare Olmo 3 to would be Qwen3, as the Qwen3 series has two models of similar size, and the Qwen3 models have a similar performance.

First, let's take a look at the smaller of the two, Olmo 3 7B.



As we can see, the Olmo 3 architecture is relatively similar to Qwen3. However, it's worth noting that this is essentially likely inspired by the Olmo 2 predecessor, not Qwen3.

Similar to Olmo 2, Olmo 3 still uses post-norm instead of pre-norm, as they found in the 2 paper that it stabilizes the training.

Interestingly, the 7B model still uses multi-head attention similar to Olmo 2. However, to make things more efficient and shrink the KV cache size, they now use sliding window attent (e.g., similar to Gemma 3).

Next, let's look at the 32B model.




Figure 45: Olmo 3 32B and Qwen3 32B side by side.

Overall, it's the same architecture but just scaled up. Also, the proportions (e.g., going from the input to the intermediate size in the feed forward layer, and so on) roughly match those in Qwen3.

My guess is the architecture was initially somewhat smaller than Qwen3 due to the smaller vocabulary, and they then scaled up the intermediate size expansion from 5x in Qwen3 to 5.4 in Olmo 3 to have a 32B model for a direct comparison.


Also, note that the 32B model uses grouped query attention.

Perhaps a last small detail is that Olmo 3 uses YaRN for context extension for the support of a context length of 64k, but only for the global (non-sliding-window-attention) layers. (YaRN is essentially a careful RoPE rescaling technique, which helps preserve model quality between long context sizes.)


In Qwen3, YaRN is optional to extend the native context from 32k tokens to 131k tokens.

If you are interested in additional architecture details, I implemented Olmo 3 from scratch in this standalone notebook [here](#).

Olmo 3 7B & 32B From Scratch



Different model types / pre-trained weights



*Only the architecture is implemented from scratch; models are not trained from scratch in this notebook, we load these pre-trained weights

Figure 46: Olmo 3 from-scratch implementation

16. DeepSeek V3.2

This article started with DeepSeek V3, which was released back in December 2024. There have been multiple DeepSeek releases back then, but I largely skipped them as they were not big flagship-model releases like DeepSeek V3 and DeepSeek R1.




Figure 47: A timeline of the DeepSeek model releases since DeepSeek V3. The main models are shown in red.

However, DeepSeek V3.2 was a really big release as it is on par with the current GPT-5 Gemini 3.0 Pro models on certain benchmarks.

The architecture is overall similar to DeepSeek V3 but they added a sparse attention mechanism to improve efficiency.




Figure 48: The DeepSeek model architecture with multi-head latent and sparse attention.

I originally planned to write a short section about DeepSeek V3.2 for this article, but it turned into a >5000 word write-up, so I moved it to a separate article, which I linked below:



A Technical Tour of the DeepSeek Models from V3 to V3.2

SEBASTIAN RASCHKA, PHD · 2025年12月3日

[Read full story →](#)

17. Mistral 3

On December 2, 2025, one day after the DeepSeek V3.2 release, the Mistral team released their new [Mistral 3](#) model suite. This includes three smaller dense models (3B, 8B, and 7B) under the Ministral 3 name, as well as their new Mistral 3 Large flagship model, which is a 675B parameter MoE (with 41B parameters active). More specifically, the Mistral 3 Large model consists of

- An MoE Language Model with 673B params and 39B active
- A 2.5B Vision Encoder

(Since this article focuses on the LLM aspects, we will ignore the vision encoder in this section. I should perhaps update my [multimodal LLMs article](#) sometime, though.)

First, it's interesting to note that it's Mistral's first MoE since Mixtral in 2023 (earlier in this article, I wrote that Mistral abandoned MoEs, and DeepSeek V3 last year ushered in an revival).

The release blog article says that all model sizes come in base, instruct, and reasoning variants, which is nice. However, their reasoning version of their 675B model is not available yet.

Another interesting tidbit is that Mistral partnered with NVIDIA here to optimize tokens, throughput on Blackwell chips, [according to their announcement](#). This is nice because it means the Ministral models will run a bit faster than comparable models on my little DC Spark (I still have to test this).

Besides the token/sec speed advantage of Mistral 3, based on quality benchmarks, their smaller models, Ministral, look on par with Qwen3. The larger flagship model is on with DeepSeek V3.1.

Since the release of Mistral 3 was just one day after DeepSeek V3.2's release, they did include any V3.2 comparisons in their article (except for the LMArena Elo score, where DeepSeek V3.2 is slightly ahead with 1423 vs 1418).

Unfortunately, it's not possible to do an apples-to-apples comparison right now, because Mistral 3 Large currently doesn't have a reasoning model, and DeepSeek V3.2 didn't show the benchmark results for their non-thinking mode, but in case you are curious, I overlaid DeepSeek V3.2-Thinking numbers (from the [DeepSeek V3.2 report](#)) with the Mistral 3 benchmark chart.




Figure 49: Mistral 3 Large benchmarks from the [Mistral 3 announcement](#), with the DeepSeek V3.2 results (from the [DeepSeek V3.2 paper](#)) overlayed on top of it.

Looking at the Mistral Large 3 Instruct model with the DeepSeek V3.2-Thinking model to it (numbers are from the DeepSeek V3.2 paper), the V3.2-Thinking model is obvious

much better. So, I am staying tuned for the Mistral 3 Large Thinking release and look forward to seeing the updated plot!

So, right now, I would say that, thanks to the optimizations, Mistral 3 Large is a great candidate for cost-effective, low-latency deployments. DeepSeek V3.2-Thinking is great if you want to maximize answer quality. Another selling point of Mistral 3 Large is that it has multimodal support as well (DeepSeek V3.2 is text-only).

By the way, my focus on DeepSeek V3.2 here in this section comes from the fact that these models were released so close to each other, within a day of each other. Plus, they have almost identical size, 671B and 673B, which makes for an interesting comparison!

Unfortunately, there is no technical report that contains more information about the model development. However, since it's an open-weight model, we do have the model weights on [Hugging Face](#) to analyze, though. So, let's take a closer look at Mistral 3 Large.

As it turns out, Mistral 3 Large is exactly the same architecture as DeepSeek V3 and V3.2. The only difference is that they increased the size of the experts by a factor of 2 while decreasing the number of experts by the same factor.




Figure 50: DeepSeek V3 and Mistral 3 Large side by side.

However, while it's effectively the same architecture, it is likely the Mistral team trained 3 from scratch rather than initializing it from DeepSeek V3 and further training it, because Mistral uses its own tokenizer.

Next to Kimi K2, Mistral 3 is now the second model series to use the DeepSeek V3 architecture. However, where the Kimi K2 team scaled up the model size from 671B to 1.3 trillion, the Mistral 3 team only changed the expert size ratio and added a vision encoder for multimodal support. But yes, why not? I think DeepSeek V3 is a pretty solid architectural design, plus it has these nice MoE and MLA efficiency aspects to it. So, why change what ain't broke? A lot of the secret sauce these days is in the training pipeline as well as the inference scaling strategies.

18. Nemotron 3 Nano and Super

This article is not an exhaustive list of all LLMs out there. To keep it manageable, I am focusing on the main highlights. Here, "highlights" means that they are either very popular, perform very well, or have an interesting architecture component.

That being said, it's time to finally add one of NVIDIA's models to this list. NVIDIA [just released](#) their newest entry in the Nemotron series, Nemotron 3, on December 15th, 2024. What's nice about Nemotron is, is that it doesn't come with just the open weights and a [technical report](#), but NVIDIA also shares the [dataset](#) and [training code](#) similar to Olmo 3.

According to the [announcement article](#), Nemotron 3 comes in three sizes:

1. Nano (30B-A3B),
2. Super (100B), (later, this was updated to 120B, see section 18.1)
3. and Ultra (500B).

18.1 Nemotron 3 Nano

Architecture-wise, the models are a Mixture-of-Experts (MoE) Mamba-Transformer hybrid architecture. As of this writing (Dec 17), only the Nano model has been released as open weight model, so the discussion below will focus on it, as illustrated in the figure below.

Nemotron 3 Nano (30B-A3B)




Figure 51.1: Outline of the Nemotron 3 Nano model, which is a Transformer-Mamba hybrid.

As illustrated above, Nemotron 3 Nano (30B-A3B) is a 52-layer hybrid Mamba-Transformer model that interleaves [Mamba-2](#) sequence-modeling blocks with sparse Mixture-of-Experts (MoE) feed-forward layers, and uses self-attention only in a small subset of layers.

There's a lot going on in the figure above, but in short, the architecture is organized into macro blocks with repeated Mamba-2 → MoE sub-blocks, plus a few Grouped-Query Attention layers. In total, if we multiply the macro- and sub-blocks, there are 52 layers in the architecture.

Regarding the MoE modules, each MoE layer contains 128 experts but activates only 1 shared and 6 routed experts per token.

The Mamba-2 layers would take a whole article itself to explain (perhaps a topic for another time). But for now, conceptually, you can think of them as similar to the Gated DeltaNet approach that Qwen3-Next and Kimi-Linear use, which I introduced above. You can also read more about it in my other Beyond Standard LLMs article:



Beyond Standard LLMs

SEBASTIAN RASCHKA, PHD • 2025年11月4日

[Read full story →](#)

The similarity between Gated DeltaNet and Mamba-2 layers is that both replace standard attention with a gated-state-space update. The idea behind this state-space-style model is that it maintains a running hidden state and mixes new inputs via learned gates. In contrast to standard attention, it scales linearly instead of quadratically with the input sequence length.

What's actually quite exciting about this architecture is its really good performance compared to pure transformer architectures of similar size, while achieving much higher tokens-per-second throughput.

Overall, this is an interesting direction, even more extreme than Qwen3-Next and Kimi-Linear in its use of only a few attention layers. However, one of the strengths of the transformer architecture is its performance at a (really) large scale. I am curious to see how Nemotron Super and especially Ultra will compare to the likes of DeepSeek V3.2.

18.2 Nemotron 3 Super

On March 11, 2026, NVIDIA now also released the 120B Super version as [open-weight models](#) on the Hugging Face Hub alongside a nice new “[Super](#)”-focused technical rep

Compared to the Nano model, besides scaling the architecture, there are two main modifications to the architecture.

First Nemotron 3 Super uses [Multi-Token Prediction \(MTP\)](#), which is a technique that t the LLM to predict multiple future tokens at each step, rather than a single one.

Instead of training the model only with the standard next-token objective, MTP also tra to predict multiple future token offsets from the same position. This provides a richer tr signal and, according to the Super report, improves both modeling quality and inferenc efficiency.




Figure 51.2: Multi-Token Prediction versus regular next token prediction. (Left subfigure inspired by the [MTP paper](#).) Originally, MTP was only used during training, not inference; hence, the inference time steps (bottom) show a single next-token prediction.

A key difference from the standard uses of MTP (which I've drawn in figure 51.2 above) that Nemotron 3 Super does not use it only during training.

The Nemotron 3 Super explicitly uses MTP at inference time as well, where the shared weight MTP head acts as an internal draft model for native speculative decoding. During generation, the model can then propose candidate continuations and then verify them with the main model. This reduces inference latency without needing a separate external draft model.

Since this is not quite standard MTP, it is perhaps more accurate to describe Nemotron 3 Super as using shared-weight MTP for speculative decoding than to call it something like "MTP-3" like in other architectures (like Step 3.5 Flash, which I covered [here](#)).

The second main difference compared to Nano is that the Super architecture uses late token experts, meaning that the experts operate in latent space (the inputs to the MoE layer are down-projected from 4096 to 1024 dimensions, the experts are applied, and then the outputs are up-projected back from 1024 to 4096 dimensions).

Nemotron 3 Super (120B-A12B)





Figure 51.3: Nemotron 3 Super 120B-A12B with latent MoE layers, multi-token prediction, and the Mamba-2 hybrid attention approach.

Benchmark-wise Nemotron 3 Super is on par with Qwen3.5 122B-A10B and GPT-OS 120B, but the throughput, thanks to the aforementioned “tricks” (MTP, latent MoE, and attention) is great: 2x faster than Qwen3.5 122B-A10B and (regarding the NVFP4 vers 2.2x faster than GPT-OSS 120B.

Nemotron 3 is 2.2x faster than GPT-OSS in NVFP4 precision



Nemotron 3 is 2x faster than Qwen3.5 in bf16 precision


Figure 51.4: Nemotron 3 Super benchmark comparison from the [Hugging Face Hub](#) page.

19. Xiaomi MiMo-V2-Flash

There's been another impressive entry in December 2025. Xiaomi released their newest LLM, Xiaomi MiMo-V2-Flash with impressive benchmark performance matching DeepSeek V3 while only having half the parameters and being faster in inference. It's a 309B Mixture of Experts (MoE) model with 15 active parameters per token.

Interestingly, it uses sliding window attention (SWA) in a 5:1 ratio with global (regular) attention, similar to Gemma 3 (see section 3). However, it uses a much more aggressive sliding window size (128) that is 8 times smaller than Gemma 3 (1024).

Xiaomi MiMo-V2-Flash (309B)



DeepSeek V3.2 (671B)




Figure 52: Xiaomi MiMo-V2-Flash compared to DeepSeek V3.2, which has similar benchmark performance.

Based on my knowledge, this is the largest sliding window attention model to date.

Additionally, the Xiaomi model uses multi-token prediction (MTP), as described earlier in section 12.3.

20. Arcee AI Trinity Large

It's been a while since the last LLM architecture addition. On January 27, Arcee AI (a company I hadn't had on my radar up to then) began releasing versions of their open-weight 400B [Trinity Large](#) LLMs on the model hub, along with two smaller variants.

Their flagship large model is a 400B param MoE (13B active params). The two smaller variants are Trinity Mini (26B with 3B active parameters) and Trinity Nano (6B with 1B active parameters).

Arcee AI Trinity Large (400B)




Figure 53: Overview of the Trinity Large architecture (based on the model hub [config file](#)).

Along with the model weights, Arcee AI also released a nice [technical report](#) with lots of details.

So, let's take a closer look at the 400B flagship model. The figure below compares it to previously discussed GLM 4.5 (section 11), which is perhaps the most similar and is also relatively small. Also, the Trinity technical report showed that the modeling performance of the Trinity Large and GLM-4.5 base models are practically identical (I assume they did compare it to more recent base models because many companies only share their fine-tuned models these days.)

Arcee AI Trinity Large (400B)

GLM-4.5 (355B)




Figure 54: Arcee AI Trinity Large next to GLM 4.5 of a relatively similar size (400B vs 355B).


But as we can see, there are several interesting architectural components added to the model.

First, there are the alternating local:global (sliding window) attention layers like in Gemini, Olmo 3, Xiaomi MiMo, etc., earlier. But instead of using the common 5:1 ratio that Gemini and Xiaomi used, they opted for a 3:1 ratio similar to Olmo 3, and a relatively large sliding window size of 4096 (also similar to Olmo 3).

In addition to QK-Norm (covered in section 2, Olmo 2), they use NoPE in the global layer (discussed NoPE in section 7, SmoLLM3).

They also have a form of gated attention. They don't have the full-blown GatedDeltaNorm (discussed in section 12) but use a similar gating as in the attention mechanism in Qwen2.5-Next.

But they modified the standard attention by adding elementwise gating to the scaled dot product before the output linear projection (as shown in the figure below), which reduces attention sinks and improves long-sequence generalization. Additionally, it also helped training stability.



```

1  def attn_trinity(x, Wq, Wk, Wv, Wo, Wg, *, h):
2      B, T, D = x.shape
3      dh = D // h
4      q = (x @ Wq.T).view(B, T, h, dh).transpose(1, 2)
5      k = (x @ Wk.T).view(B, T, h, dh).transpose(1, 2)
6      v = (x @ Wv.T).view(B, T, h, dh).transpose(1, 2)
7
8      q = rmsnorm(q)  # QK-Norm
9      k = rmsnorm(k)
10
11     o = sdpa(q, k, v) # standard scaled dot product attention
12
13     ### Gating ###
14     g = torch.sigmoid(x @ Wg.T).view(B, T, h, dh).transpose(1, 2)
15     o = o * g # Elementwise gate AFTER attention before Wo output projection
16     #####
17
18     return (o.transpose(1, 2).reshape(B, T, D)) @ Wo.T

```

Figure 55: Illustration of the gating mechanism that Trinity Large uses in the attention mechanism.

You may have noticed the use of four (instead of two) RMSNorm layers in the previous Large architecture figure. This is their so-called depth-scaled sandwich norm, which is on previous work but something I haven't seen before in major architectures. Overall, it like a Gemma 3-like RMSNorm placement, but the twist here is that the gain of the sec RMSNorm (in each block) is depth-scaled, meaning it's initialized to about $1 / \sqrt{L}$ (L the total number of layers). So, early in training, the residual update starts small and grc the model learns the right scale.

The MoE is a DeepSeek-like MoE with lots of small experts, but made it coarser as that with inference throughput (something we have also seen in Mistral 3 Large when they adopted the DeepSeek V3 architecture).

Lastly, there are some interesting details on the training improvements (a new MoE load balancing strategy and another using the MuOpt optimizer), but since this is an archite post, these are out of scope.

21. GLM-5

Chinese New Year has become a surprisingly reliable window for strong open-weight releases. For example, GLM-4 and Qwen 1.5 were released in January and February 2025 and DeepSeek R1 and Qwen 2.5 were released in 2025.

This year, z.AI (Zhipu AI) kicked things off (again relatively early) with GLM-5 on February 2026, approximately a week before the Lunar New Year on February 17.

Compared to the GLM-4.5 model I covered earlier in this article (see section 11, released summer 2025), its GLM-5 successor is twice the size: up from 355B parameters to 744B pushing it into the territory between DeepSeek-V3.2 and Kimi K2.

Similar to GLM-4.5, GLM-5 is a Mixture-of-Experts (section 1.2) model, and the number of active parameters per token is only increased slightly: 40B in GLM-5 versus 32B in GLM-4.5.




Figure 56: Architecture of GLM-5 and GLM-4.5 side by side.

Interestingly, as shown in Figure 56 above, GLM-5 adopts DeepSeek's multi-head late attention (MLA, see section 1.1) as well as DeepSeek Sparse Attention (which I covered

more detail in my [DeepSeek V3.2](#) article). The motivation for these modifications is to reduce the inference cost when working with long contexts.

Other than that, the architecture is relatively similar. The increased size is mainly due to having more experts (256 instead of 160) and slightly increasing the layer sizes. For instance, the embedding dimension and expert size are now 6,144 (up from 5,120), an intermediate projection size is also slightly up from 1,536 to 2,048. Interestingly, the number of layers (transformer blocks) is reduced from 92x to 78x. I am assuming this is to reduce inference costs and make the model faster (because layer depth can't be parallelized).

I usually don't include benchmarks here since this article is focused on the architecture. If we were to include training details and evaluations, this article would grow way out of scope and length. That being said, I saw that I included the GLM-4.5 benchmark back in July 2023, so I will make another exception here, because the benchmarks look truly impressive and competitive with all major flagship LLM offerings (GPT-5.2 extra-high, Gemini Pro 3, and Claude 4. Opus). But again, it's worth highlighting that benchmark performance isn't necessarily indicative of real-world performance.




Figure 57: GLM architectures next to benchmarks. The GLM-4.7 architecture is similar to GLM-4.5. The benchmarks are taken from the GLM-5 release blog post: <https://z.ai/blog/glm-5>

22. More February 2026 Releases: from Kimi K2.5 to Tiny Aya


In total, there have been 10 interesting open-weight LLM releases between January and February 2026:

1. Arcee AI's Trinity Large (Jan 27, 2026)
2. Moonshot AI's Kimi K2.5 (Jan 27, 2026)
3. StepFun Step 3.5 Flash (Feb 1, 2026)
4. Qwen3-Coder-Next (Feb 3, 2026)
5. z.AI's GLM-5 (Feb 12, 2026)
6. MiniMax M2.5 (Feb 12, 2026)
7. Nanbeige 4.1 3B (Feb 13, 2026)
8. Qwen 3.5 (Feb 15, 2026)
9. Ant Group's Ling 2.5 1T & Ring 2.5 1T (Feb 16, 2026)
10. Cohere's Tiny Aya (Feb 17, 2026)
11. Sarvam 30B and 105B (Mar 6, 2026)

I covered Arcee AI's Trinity Large and z.AI's GLM-5 in sections 19 and 20 above. However, since there was a lot of content to cover for the January-February time period, I wrote a standalone article with more information about the 10 architectures listed above here:

A Dream of Spring for Open-Weight LLMs: 10 Architectures from Jan-Feb 2026

SEBASTIAN RASCHKA, PHD • 2月25日



If you have struggled a bit to keep up with open-weight model releases this month, this article should c
you up on the main themes.

[Read full story →](#)

23. Gemma 4

After the Nemotron 3 Super release in March, the rest of the month was relatively quiet
flagship open-weight model releases. While I am still waiting for DeepSeek-V4, April at
brought us Google's Gemma 4.

Architecture-wise, Gemma 4 (31B) looks pretty much unchanged compared to Gemm
(27B), as illustrated in the figure below.




Figure 58: Gemma 3 (27B) and Gemma 4 (31B) side by side.

(Note that Gemma 4 also has multimodal model support now, but I will leave the image encoder part for a separate article in the future; here, we only focus on the text portion.)

As we can see in the figure above, Gemma 4 maintains a relatively unique Pre- and Post-norm setup and remains relatively classic, with a 5:1 hybrid attention mechanism combining a sliding-window (local) layer and a full-attention (global) layer. The attention mechanism itself is also classic Grouped Query Attention (GQA).

However, a small change over Gemma 3, which is easy to overlook, is that for the global attention layers they reuse the keys in the attention mechanism. I.e., they set values = keys which should result in further KV cache size reduction.

Furthermore, Gemma 4 also uses p-RoPE, where only 25% of the frequency pairs get positional information. This helps with reducing positional noise in long-context situations.

But let's not be fooled by the lack of big(ger) architectural changes. Looking at the benchmarks, Gemma 4 is a huge leap from Gemma 3! For instance, on the [AI Arena Leaderboard](#), Gemma 4 (31B) ranks similarly to the much larger Qwen3.5-397B-A17B model. But as I discussed in my model evaluation article (linked below), arena scores are problematic as they can be gamed and are biased towards human (style) preference.



Understanding the 4 Main Approaches to LLM Evaluation (From Scratch)

SEBASTIAN RASCHKA, PHD · 2025年10月5日

[Read full story →](#)

However, if we look at some other common benchmarks, which I plotted below, we can see that it's indeed a very clear leap over Gemma 3 and ranks on par with Qwen3.5 27B.




Figure 59: Gemma 3 versus Gemma 4 versus Qwen3.5 (the numbers are taken from the [Gemma 4](#) and [Qwen3.5](#) model hub pages).

Note that there is also a Mixture-of-Experts (MoE) Gemma 4 variant, which is illustrated below next to a Qwen3 model of similar size.




Figure 60: Qwen3 Coder Flash compared to Gemma 4 MoE.

As the figure above shows, the approaches are relatively similar except that Gemma 4 the unique Pre- and Post-norm placement discussed earlier.

Benchmark-wise, the Gemma 4 MoE variant, which has 4B parameters less in total than the Gemma 4 (31B) dense variant, the performances are relatively similar.




Figure 61: Gemma 4 MoE (26B-A4B) is only slightly worse than Gemma 4 (31) dense.

If you are interested in a visual overview of all the architectures covered here, I put together an LLM Architecture Gallery [here](#).






LLM architecture gallery at <https://sebastianraschka.com/llm-architecture-gallery/>

After all these years, LLM releases remain exciting, and I am curious to see what next!

This magazine is a personal passion project, and your support helps keep it alive.

If you'd like to support my work, please consider my [Build a Large Language Model \(From Scratch\)](#) book or its follow-up, [Build a Reasoning Model \(From Scratch\)](#). (I'm confident you'll get a lot out of these; they explain how LLMs work in depth you won't find elsewhere.)

Thanks for reading, and for helping support independent research!



Build a Large Language Model (From Scratch) is now available on [Amazon](#). Build a Reasoning Model (From Scratch) is in [Early Access at Manning](#).

If you read the book and have a few minutes to spare, I'd really appreciate a [brief review](#) helps us authors a lot!

Your support means a great deal! Thank you!

Subscribe to Ahead of AI

By Sebastian Raschka · Thousands of paid subscribers

Ahead of AI focuses on machine learning and AI research and is read by more than 190,0 researchers and practitioners who want to stay ahead in a rapidly evolving field.

Upgrade to paid

By subscribing, you agree Substack's [Terms of Use](#), and acknowledge its [Information Collection Notice](#) and [Privacy Policy](#).



1,975 Likes · 174 Restacks

Discussion about this post

[Comments](#) [Restacks](#)

Write a comment...



Daniel Kleine 2025年7月21日

♥ Liked by Sebastian Raschka, PhD

Apart from the architectural differences, what would be interesting to know is on which data the LLMs have been trained on. From my pov, it's really unfortunate that this info is typically not disclosed, even for open-source LLMs. Not just the amount of training data (number of tokens) but also the data quality as factors for a true scientific comparison.

♥ LIKE (18) 💬 REPLY

2 replies by Sebastian Raschka, PhD and others



Leo Benaharon 2025年7月19日

♥ Liked by Sebastian Raschka, PhD

Amazing article! This is evidence that we haven't hit a wall yet with LLMs as all these LLMs haven't converged to the same architectures.

Cohere Labs is also doing some great work for open source and have some interesting insights. I feel a lot of people don't know who they are as they are trying to appeal to businesses/governments.

♥ LIKE (9) 💬 REPLY

1 reply by Sebastian Raschka, PhD

95 more comments...

© 2026 Raschka AI Research (RAIR) Lab LLC · [Privacy](#) · [Terms](#) · [Collection notice](#)
[Substack](#) is the home for great culture